

# Session-5

## Advance Hive & HBASE BASICS

### Assignment

---

#### Task 1.1

1. Write a Hive program to find the number of medals won by each country in swimming.

```
select country, sum(total_medals) as medals from olympix_data where
sport='Swimming' group by country;
```

```
hive> select country, sum(total_medals) from olympix_data where sport='Swimming' group by country;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a diff
Query ID = acadgild_20181118182852_30b7ef9a-56f4-41ae-81f1-42c3409c02ce
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1542542661295_0003, Tracking URL = http://localhost:8088/proxy/application_1542542661295_0003/
Kill Command = /home/acadgild/install/hadoop/hadoop-2.6.5/bin/hadoop job -kill job_1542542661295_0003
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2018-11-18 18:29:13,786 Stage-1 map = 0%, reduce = 0%
2018-11-18 18:29:30,885 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 5.34 sec
2018-11-18 18:29:48,808 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 9.95 sec
MapReduce Total cumulative CPU time: 9 seconds 950 msec
Ended Job = job_1542542661295_0003
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 9.95 sec HDFS Read: 528693 HDFS Write: 881 SUCCESS
Total MapReduce CPU Time Spent: 9 seconds 950 msec
OK
Argentina      1
Australia      163
Austria        3
Belarus        2
Brazil         8
Canada         5
China          35
Costa Rica     2
Croatia        1
Denmark        1
France         39
Germany        32
Great Britain  11
Hungary        9
Italy          16
Japan          43
Lithuania      1
Netherlands    46
Norway         2
Poland         3
Romania        6
Russia         20
Serbia         1
Slovakia       2
Slovenia       1
South Africa   11
South Korea    4
Spain          3
Sweden         9
Trinidad and Tobago 1
```

2. Write a Hive program to find the number of medals that India won year wise.

```
select year, sum(total_medals) as medals from olympix_data where  
country='India' group by year;
```

```
Time taken: 63.084 seconds, Fetched: 4 row(s)  
[hive> select year, sum(total_medals) from olympix_data where country='India' group by year;  
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a diff  
Query ID = acadgild_20181118183533_3e243020-3ce9-4ec3-af04-fcf6563ea651  
Total jobs = 1  
Launching Job 1 out of 1  
Number of reduce tasks not specified. Estimated from input data size: 1  
In order to change the average load for a reducer (in bytes):  
  set hive.exec.reducers.bytes.per.reducer=<number>  
In order to limit the maximum number of reducers:  
  set hive.exec.reducers.max=<number>  
In order to set a constant number of reducers:  
  set mapreduce.job.reduces=<number>  
Starting Job = job_1542542661295_0004, Tracking URL = http://localhost:8088/proxy/application_1542542661295_0004/  
Kill Command = /home/acadgild/install/hadoop/hadoop-2.6.5/bin/hadoop job -kill job_1542542661295_0004  
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1  
2018-11-18 18:35:54,569 Stage-1 map = 0%, reduce = 0%  
2018-11-18 18:36:16,240 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 5.57 sec  
2018-11-18 18:36:34,048 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 9.29 sec  
MapReduce Total cumulative CPU time: 9 seconds 290 msec  
Ended Job = job_1542542661295_0004  
MapReduce Jobs Launched:  
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 9.29 sec HDFS Read: 528698 HDFS Write: 163 SUCCESS  
Total MapReduce CPU Time Spent: 9 seconds 290 msec  
OK  
2000 1  
2004 1  
2008 3  
2012 6  
Time taken: 63.084 seconds, Fetched: 4 row(s)
```

3. Write a Hive Program to find the total number of medals each country won.

```
select country, sum(total_medals) as medals from olympix_data group by
country order by medals desc;
```

```
[hive> select country, sum(total_medals) as medals from olympix_data group by country order by medals desc;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different
Query ID = acadgild_20181118184048_7e12911f-4809-4fba-ba26-e75f1b1faf70
Total jobs = 2
Launching Job 1 out of 2
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1542542661295_0005, Tracking URL = http://localhost:8088/proxy/application_1542542661295_0005/
Kill Command = /home/acadgild/install/hadoop/hadoop-2.6.5/bin/hadoop job -kill job_1542542661295_0005
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2018-11-18 18:41:10,390 Stage-1 map = 0%, reduce = 0%
2018-11-18 18:41:28,178 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 4.59 sec
2018-11-18 18:41:45,318 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 8.72 sec
MapReduce Total cumulative CPU time: 8 seconds 720 msec
Ended Job = job_1542542661295_0005
Launching Job 2 out of 2
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1542542661295_0006, Tracking URL = http://localhost:8088/proxy/application_1542542661295_0006/
Kill Command = /home/acadgild/install/hadoop/hadoop-2.6.5/bin/hadoop job -kill job_1542542661295_0006
Hadoop job information for Stage-2: number of mappers: 1; number of reducers: 1
2018-11-18 18:42:13,578 Stage-2 map = 0%, reduce = 0%
2018-11-18 18:42:29,895 Stage-2 map = 100%, reduce = 0%, Cumulative CPU 3.18 sec
2018-11-18 18:42:49,040 Stage-2 map = 100%, reduce = 100%, Cumulative CPU 7.12 sec
MapReduce Total cumulative CPU time: 7 seconds 120 msec
Ended Job = job_1542542661295_0006
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 8.72 sec HDFS Read: 527007 HDFS Write: 3132 SUCCESS
Stage-Stage-2: Map: 1 Reduce: 1 Cumulative CPU: 7.12 sec HDFS Read: 8693 HDFS Write: 2742 SUCCESS
Total MapReduce CPU Time Spent: 15 seconds 840 msec
OK
United States 1312
Russia 768
Germany 629
Australia 609
China 530
Canada 370
Italy 331
Great Britain 322
France 318
Netherlands 318
South Korea 308
Japan 282
Brazil 221
```

4. Write a Hive program to find the number of gold medals each country won.

```
select country, sum(golds) as gold_medals from olympix_data group by
country order by gold_medals desc;
```

```
[hive> select country, sum(golds) as gold_medals from olympix_data group by country order by gold_medals desc;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different
Query ID = acadgild_20181118184956_bf585796-b4c7-4c2c-bff1-a64de39d8a67
Total jobs = 2
Launching Job 1 out of 2
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1542542661295_0007, Tracking URL = http://localhost:8088/proxy/application_1542542661295_0007/
Kill Command = /home/acadgild/install/hadoop/hadoop-2.6.5/bin/hadoop job -kill job_1542542661295_0007
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2018-11-18 18:50:16,714 Stage-1 map = 0%, reduce = 0%
2018-11-18 18:50:38,177 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 3.47 sec
2018-11-18 18:50:55,007 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 7.37 sec
MapReduce Total cumulative CPU time: 7 seconds 370 msec
Ended Job = job_1542542661295_0007
Launching Job 2 out of 2
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1542542661295_0008, Tracking URL = http://localhost:8088/proxy/application_1542542661295_0008/
Kill Command = /home/acadgild/install/hadoop/hadoop-2.6.5/bin/hadoop job -kill job_1542542661295_0008
Hadoop job information for Stage-2: number of mappers: 1; number of reducers: 1
2018-11-18 18:51:21,864 Stage-2 map = 0%, reduce = 0%
2018-11-18 18:51:38,318 Stage-2 map = 100%, reduce = 0%, Cumulative CPU 2.87 sec
2018-11-18 18:51:57,805 Stage-2 map = 100%, reduce = 100%, Cumulative CPU 6.43 sec
MapReduce Total cumulative CPU time: 6 seconds 430 msec
Ended Job = job_1542542661295_0008
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 7.37 sec HDFS Read: 526989 HDFS Write: 3107 SUCCESS
Stage-Stage-2: Map: 1 Reduce: 1 Cumulative CPU: 6.43 sec HDFS Read: 8660 HDFS Write: 2703 SUCCESS
Total MapReduce CPU Time Spent: 13 seconds 800 msec
OK
United States 552
China 234
Russia 234
Germany 223
Canada 168
Australia 163
Great Britain 124
South Korea 110
France 108
Netherlands 101
Norway 97
Italy 86
Hungary 77
```

### Task 1.2

Write a hive UDF that implements functionality of string concat\_ws(string SEP, array<string>).

This UDF will accept two arguments, one string and one array of string. It will return a single string where all the elements of the array are separated by the SEP.

```
1 package com.acadgild.custom_udf;
2
3 import java.util.List;
4
5 import org.apache.hadoop.hive.ql.exec.UDF;
6 import org.apache.hadoop.io.Text;
7
8 public class ConcatWS extends UDF {
9     public Text evaluate(Text sep, List<String> list) {
10         String result = "";
11         Text tx = new Text();
12         for(String str : list) {
13             result = result.concat(str.concat(sep.toString()));
14         }
15         tx.set(result.substring(0, result.length()-1));
16         return tx;
17     }
18 }
19
```

### Task 1.3

#### ACID Transactions – Transactions in Hive

```
hive> update temperature_data set temperature=25 where zip_code = 560037;
```

FAILED: SemanticException [Error 10294]: Attempt to do update or delete using transaction manager that does not support these operations.

```
create table temperature_orc(temp_data date, zip_code int, temperature int) clustered by (zip_code) into 3 buckets stored as orc TBLPROPERTIES('transactional'='true');
```

```
[hive> insert into temperature_orc select * from temperature_data;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using the newer Hive MR flavor.
Query ID = acadgild_20181118214626_1a122db6-ad43-44e5-91eb-f657f1b1873d
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks determined at compile time: 3
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1542542661295_0010, Tracking URL = http://localhost:8088/proxy/application_1542542661295_0010
Kill Command = /home/acadgild/install/hadoop/hadoop-2.6.5/bin/hadoop job -kill job_1542542661295_0010
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 3
2018-11-18 21:46:45,732 Stage-1 map = 0%, reduce = 0%
2018-11-18 21:46:59,896 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 2.72 sec
2018-11-18 21:47:33,289 Stage-1 map = 100%, reduce = 22%, Cumulative CPU 4.99 sec
2018-11-18 21:47:34,644 Stage-1 map = 100%, reduce = 44%, Cumulative CPU 6.7 sec
2018-11-18 21:47:38,716 Stage-1 map = 100%, reduce = 67%, Cumulative CPU 10.13 sec
2018-11-18 21:47:40,252 Stage-1 map = 100%, reduce = 78%, Cumulative CPU 12.33 sec
2018-11-18 21:47:42,938 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 15.5 sec
MapReduce Total cumulative CPU time: 15 seconds 500 msec
Ended Job = job_1542542661295_0010
Loading data to table custom.temperature_orc
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 3 Cumulative CPU: 15.5 sec HDFS Read: 17215 HDFS Write: 2461 SUCCESS
Total MapReduce CPU Time Spent: 15 seconds 500 msec
OK
Time taken: 79.994 seconds
```

### INSERT

```
insert into temperature_orc values ('2018-11-18',560037,22);
```

```
[hive> select * from temperature_orc;
OK
2018-11-18      560037      22
1993-01-10      123112      11
1991-01-10      123112      11
1991-01-10      123112      11
1990-01-10      123112      10
1990-03-10      381920      15
1991-02-12      384902      10
1994-01-10      302918      23
1993-03-10      381920      16
1994-02-14      283901      12
1991-02-12      384902      10
1991-02-12      384902      10
1990-01-10      302918      23
1991-03-10      381920      16
1990-02-14      283901      12
1991-02-14      283901      11
1990-02-12      384902      9
1991-01-10      302918      22
1990-01-10      302918      23
1991-03-10      381920      16
1990-02-14      283901      12
Time taken: 0.42 seconds, Fetched: 21 row(s)
```



```
[hive> insert into temperature_orc values ('2018-11-18',560037,22);
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future ve
Query ID = acadgild_20181118220358_a391867f-d164-416a-8623-b34345b63f19
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks determined at compile time: 3
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1542542661295_0012, Tracking URL = http://localhost:8088/proxy/app
Kill Command = /home/acadgild/install/hadoop/hadoop-2.6.5/bin/hadoop job -kill job_1
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 3
2018-11-18 22:04:18,053 Stage-1 map = 0%, reduce = 0%
2018-11-18 22:04:34,804 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 5.11 sec
2018-11-18 22:05:05,409 Stage-1 map = 100%, reduce = 22%, Cumulative CPU 6.96 sec
2018-11-18 22:05:09,515 Stage-1 map = 100%, reduce = 44%, Cumulative CPU 9.61 sec
2018-11-18 22:05:13,655 Stage-1 map = 100%, reduce = 67%, Cumulative CPU 13.16 sec
2018-11-18 22:05:20,519 Stage-1 map = 100%, reduce = 78%, Cumulative CPU 16.68 sec
2018-11-18 22:05:24,518 Stage-1 map = 100%, reduce = 89%, Cumulative CPU 20.65 sec
2018-11-18 22:05:25,592 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 24.04 sec
MapReduce Total cumulative CPU time: 24 seconds 40 msec
Ended Job = job_1542542661295_0012
Loading data to table custom.temperature_orc
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 3 Cumulative CPU: 24.04 sec HDFS Read: 18210 HDFS
Total MapReduce CPU Time Spent: 24 seconds 40 msec
OK
Time taken: 90.109 seconds
```

## UPDATE

update temperature\_orc set temperature=25 where zip\_code = 560037;

```
[hive> update temperature_orc set temperature=25 where zip_code = 560037;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider us
Query ID = acadgild_20181118221134_c281e974-5099-47e6-870f-b3bd2e51ce2d
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks determined at compile time: 3
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1542542661295_0013, Tracking URL = http://localhost:8088/proxy/application_1542542661
Kill Command = /home/acadgild/install/hadoop/hadoop-2.6.5/bin/hadoop job -kill job_1542542661295_0013
Hadoop job information for Stage-1: number of mappers: 3; number of reducers: 3
2018-11-18 22:11:54,391 Stage-1 map = 0%, reduce = 0%
2018-11-18 22:12:45,203 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 17.23 sec
2018-11-18 22:13:24,000 Stage-1 map = 100%, reduce = 22%, Cumulative CPU 18.94 sec
2018-11-18 22:13:26,815 Stage-1 map = 100%, reduce = 67%, Cumulative CPU 23.89 sec
2018-11-18 22:13:33,733 Stage-1 map = 100%, reduce = 89%, Cumulative CPU 28.43 sec
2018-11-18 22:13:35,042 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 30.53 sec
MapReduce Total cumulative CPU time: 30 seconds 530 msec
Ended Job = job_1542542661295_0013
Loading data to table custom.temperature_orc
MapReduce Jobs Launched:
Stage-Stage-1: Map: 3 Reduce: 3 Cumulative CPU: 30.53 sec HDFS Read: 33508 HDFS Write: 893 SUCCESS
Total MapReduce CPU Time Spent: 30 seconds 530 msec
OK
Time taken: 124.413 seconds
```

```
[hive> select * from temperature_orc;
OK
2018-11-18      560037  25
1993-01-10      123112  11
1991-01-10      123112  11
1991-01-10      123112  11
1990-01-10      123112  10
1990-03-10      381920  15
1991-02-12      384902  10
1994-01-10      302918  23
1993-03-10      381920  16
1994-02-14      283901  12
1991-02-12      384902  10
1991-02-12      384902  10
1990-01-10      302918  23
1991-03-10      381920  16
1990-02-14      283901  12
1991-02-14      283901  11
1990-02-12      384902   9
1991-01-10      302918  22
1990-01-10      302918  23
1991-03-10      381920  16
1990-02-14      283901  12
Time taken: 0.641 seconds, Fetched: 21 row(s)
```

```
hive> update temperature_orc set zip_code=560035 where zip_code =
560037;
FAILED: SemanticException [Error 10302]: Updating values of bucketing
columns is not supported. Column zip_code.
hive>
```

## **DELETE**

```
delete from temperature_orc where zip_code=381920;
```

```
[hive> delete from temperature_orc where zip_code=381920;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions
Query ID = acadgild_20181118221728_dbcbb9c5-12bb-470e-a556-c3b834ed6951
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks determined at compile time: 3
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1542542661295_0014, Tracking URL = http://localhost:8088/proxy/application_1542542661295_0014/
Kill Command = /home/acadgild/install/hadoop/hadoop-2.6.5/bin/hadoop job -kill job_1542542661295_0014
Hadoop job information for Stage-1: number of mappers: 3; number of reducers: 3
2018-11-18 22:17:48,772 Stage-1 map = 0%, reduce = 0%
2018-11-18 22:18:35,441 Stage-1 map = 67%, reduce = 0%, Cumulative CPU 14.52 sec
2018-11-18 22:18:36,516 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 16.77 sec
2018-11-18 22:19:14,224 Stage-1 map = 100%, reduce = 67%, Cumulative CPU 21.85 sec
2018-11-18 22:19:22,666 Stage-1 map = 100%, reduce = 89%, Cumulative CPU 27.62 sec
2018-11-18 22:19:24,051 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 29.89 sec
MapReduce Total cumulative CPU time: 29 seconds 890 msec
Ended Job = job_1542542661295_0014
Loading data to table custom.temperature_orc
MapReduce Jobs Launched:
Stage-Stage-1: Map: 3 Reduce: 3 Cumulative CPU: 29.89 sec HDFS Read: 33112 HDFS Write: 1024
Total MapReduce CPU Time Spent: 29 seconds 890 msec
OK
Time taken: 119.566 seconds
```



```

Time taken: 117.388 seconds
[hive> select * from temperature_orc;
OK
2018-11-18      560037  25
1993-01-10      123112  11
1991-01-10      123112  11
1991-01-10      123112  11
1990-01-10      123112  10
1991-02-12      384902  10
1994-01-10      302918  23
1994-02-14      283901  12
1991-02-12      384902  10
1991-02-12      384902  10
1990-01-10      302918  23
1990-02-14      283901  12
1991-02-14      283901  11
1990-02-12      384902  9
1991-01-10      302918  22
1990-01-10      302918  23
1990-02-14      283901  12
Time taken: 0.422 seconds, Fetched: 17 row(s)

```

## Task 2.1

1.What is NoSQL data base?

NoSQL data base is a database management technology which does not strictly follow all the rules (ACID properties) of relational DBMS and data can not be queried using traditional sql language. That is why it is termed as NO-SQL or Not Only SQL.

2.How does data get stored in NoSql database?

Data can be stored in four different manners :

**Columnar Databases** – Reads and writes columns of data rather than the rows. Each column is comparable to a container in RDBMS where a Key defines a row and single row has multiple columns.

**Document Databases** – Store and retrieve data in semi-structured format such as XML, JSON, etc.

**Graph Databases** – Stores data as entities and relations between them allowing faster traversal and joining operations to be performed.

**In-Memory Key-Value Stores**– Stores critical data in memory which in turn improves the performance of the systems.

3.What is a column family in HBase?

In the HBase data model columns are grouped into column families, which must be defined during table creation. A column family defines shared features to all columns that are created within them.

4.How many maximum number of columns can be added to HBase table?

There is no limit of number of column qualifiers.

5.Why columns are not defined at the time of table creation in HBase?

column-qualifier is not defined during table creation. This is what makes HBase schema-less. Once table is created a row can be added which should belong to at-least one column-family in a table and then can belong to any column-qualifier. If the column-qualifier is not already present, it will be created.

6. How does data get managed in HBase?

Just like in a Relational Database, data in HBase is stored in Tables and these Tables are stored in Regions. When a Table becomes too big, the Table is partitioned into multiple Regions. These Regions are assigned to Region Servers across the cluster.

7. What happens internally when new data gets inserted into HBase table?

When the client issues a Put request:

1. write the data to the write-ahead log.
2. Once the data is written to the WAL, it is placed in the MemStore. Then, the put request acknowledgement returns to the client.
3. The MemStore stores updates in memory as sorted KeyValues, the same as it would be stored in an HFile. There is one MemStore per column family.
4. When the MemStore accumulates enough data, the entire sorted set is written to a new HFile in HDFS.

## Task 2.2

1. Create an HBase table named 'clicks' with a column family 'hits' such that it should be able to store last 5 values of qualifiers inside 'hits' column family.

```
hbase(main):002:0> create 'clicks','hits'
0 row(s) in 1.5950 seconds

=> Hbase::Table - clicks
hbase(main):003:0> put 'clicks','blog','hits:visitor','Chris Lawther'
0 row(s) in 0.6160 seconds

hbase(main):004:0> put 'clicks','blog','hits:pageView','20'
0 row(s) in 0.0490 seconds

hbase(main):005:0> put 'clicks','blog','hits:dailyUnique','5'
0 row(s) in 0.0530 seconds

hbase(main):006:0> put 'clicks','blog','hits:weeklyUnique','9'
0 row(s) in 0.0240 seconds

hbase(main):007:0> put 'clicks','blog','hits:monthlyUnique','12'
0 row(s) in 0.0250 seconds
```

```
hbase(main):008:0> get 'clicks','blog'
COLUMN                                CELL
hits:dailyUnique                      timestamp=1542654342280, value=5
hits:monthlyUnique                   timestamp=1542654369622, value=12
hits:pageView                        timestamp=1542654319772, value=20
hits:visitor                         timestamp=1542654287560, value=Chris Lawther
hits:weeklyUnique                    timestamp=1542654356730, value=9
5 row(s) in 0.1400 seconds
```

5. Add few records in the table and update some of them. Use IP Address as row-key. Scan the table to view if all the previous versions are getting displayed.

```
[hbase(main):021:0> get 'clicks','192.168.0.115'
COLUMN                                CELL
hits:dailyUnique                      timestamp=1542655042473, value=12
hits:monthlyUnique                   timestamp=1542655077802, value=66
hits:pageView                        timestamp=1542655029360, value=22
hits:type                           timestamp=1542655008295, value=Blog
hits:visitor                         timestamp=1542654958931, value=Chris Lawther
hits:weeklyUnique                    timestamp=1542655059932, value=15
6 row(s) in 0.0780 seconds
```

```
[hbase(main):022:0> alter 'clicks',NAME=>'visitor',VERSIONS=>5
Updating all regions with the new schema...
0/1 regions updated.
1/1 regions updated.
Done.
0 row(s) in 3.2740 seconds

[hbase(main):023:0> put 'clicks','192.168.0.115','hits:vistor','Nick Chiota'
0 row(s) in 0.0270 seconds

[hbase(main):024:0> put 'clicks','192.168.0.115','hits:vistor','Bill fuller'
0 row(s) in 0.0210 seconds

[hbase(main):025:0> put 'clicks','192.168.0.115','hits:vistor','Anthony Passqrute'
0 row(s) in 0.0130 seconds

[hbase(main):026:0> scan 'clicks'
ROW                                     COLUMN+CELL
192.168.0.115                         column=hits:dailyUnique, timestamp=1542655042473, value=12
192.168.0.115                         column=hits:monthlyUnique, timestamp=1542655077802, value=66
192.168.0.115                         column=hits:pageView, timestamp=1542655029360, value=22
192.168.0.115                         column=hits:type, timestamp=1542655008295, value=Blog
192.168.0.115                         column=hits:visitor, timestamp=1542654958931, value=Chris Lawther
192.168.0.115                         column=hits:vistor, timestamp=1542655448152, value=Anthony Passqrute
192.168.0.115                         column=hits:weeklyUnique, timestamp=1542655059932, value=15
1 row(s) in 0.1410 seconds

[hbase(main):027:0> scan 'clicks',{COLUMN=>'hits:vistor',VERSIONS=>2}
ROW                                     COLUMN+CELL
192.168.0.115                         column=hits:vistor, timestamp=1542655448152, value=Anthony Passqrute
1 row(s) in 0.0360 seconds

[hbase(main):028:0> scan 'clicks',{COLUMN=>'hits:vistor',VERSIONS=>3}
ROW                                     COLUMN+CELL
192.168.0.115                         column=hits:vistor, timestamp=1542655448152, value=Anthony Passqrute
1 row(s) in 0.0670 seconds

[hbase(main):029:0> scan 'clicks',{COLUMN=>'hits:vistor',VERSIONS=>1}
ROW                                     COLUMN+CELL
192.168.0.115                         column=hits:vistor, timestamp=1542655448152, value=Anthony Passqrute
1 row(s) in 0.0260 seconds

[hbase(main):030:0> scan 'clicks',{COLUMN=>'hits:vistor',VERSIONS=>4}
ROW                                     COLUMN+CELL
192.168.0.115                         column=hits:vistor, timestamp=1542655448152, value=Anthony Passqrute
1 row(s) in 0.0250 seconds

hbase(main):031:0> █
```